





Predicting WHICH model gets blamed, on wrong guesses only, GroupKFold by prompt. Both feature models beat shuffled labels ($p < 0.005$) - but only because the features encode authorship. Neither beats a baseline that knows the true author and nothing else (dashed line), so the features add no information about blame beyond identity. The claim is "no signal beyond author identity", not "no signal".