


What judges say they used



Keyword matching against a hand-built lexicon (src/analyze_reasons.py) - a coarse instrument, and the categories are ours, not the judges'. A rationale can match several cues or none. Under single-text querying 96.7% of rationales name a model explicitly, against 56.9% in the lineup: judges there are reasoning from model stereotypes, which is the same collapse Fig. B2 shows in their answers. For contrast, length proxies carry 63% of Engine A's feature importance